Lab Work 2

André Pedrosa [85098], Filipe Pires [85122], João Alegria [85048]

Algorithmic Information Theory
Department of Electronics, Telecommunications and Informatics
University of Aveiro

October 22, 2019

Introduction

This report aims to describe the work developed for the first assignment of the discipline of 'Algorithmic Information Theory', explaining all the steps and decisions taken by us, and presenting the results we considered most relevant.

The programs implemented in C++ have the purpose of collecting statistical information about texts using Markov (finite-context) models, and of automatically producing texts that follow the models built.

Along with the description of the solution, we also discuss the effects of the variation of the programs' parameters and attempt to compare different types of texts by the amount of information they hold on average.

All code developed is publicly accessible in our GitHub repository:

`https://github.com/joao-alegria/TAI`

1. Information Model

Our first goal was to be able to predict the next outcome of a text source. To do this, we needed to take in consideration the dependencies between the characters of a text. The use of Markov models for the extraction of the statistical properties of a text was due to its value as an approach to represent these data dependencies.

The specific type of model that most interests us is called discrete time Markov chain or finite-context model. This model assigns a probability estimate to the symbols of the alphabet, according to a conditioning context computed over a finite and fixed number of past outcomes. More about this is explained in the description of this work assignment (*1*), containing the mathematical equations that served as the basis for our implementation.

1.1. Collecting Data

We decided to organize the program by several files, each with a different purpose, for good readability and to allow future modifications without the need of much refactoring - so we adopted a modular architecture strategy.

The file `fcm.cpp` serves as the base code for the command to be executed in order to generate a finite-context model, given one or several information source(s). By executing this command, the program is started and begins to call a set of functions that return an instance of our implementation of the Markov's model trained from the information source(s), and calculate the text entropy, as estimated by this instance. The `fcm` command has the following format:

```
$ ./fcm [-h] k alpha trainFile [trainFile ...]
```

Here, `-h` is the option that presents the manual for the command usage. Argument `k` is the value given for the size of the context. This context corresponds to a string with `k` characters, it is based on the several contexts produced and on the single characters that follow each one of them that the model is able to calculate the text's entropy. `alpha` stands for the value of the 'smoothing' parameter for estimating the probabilities of events. These events correspond to the occurrences of a character after a given context. And `trainFile` is, as the name indicates, the name of the file(s) that contain the text to be processed by the model.

It is important to mention that the `.cpp` files are never actually executed. The way the `fmc` command works is: first the user must compile the required files and then execute the generated file passing the desired arguments. We make this task easier through the given `Makefile`. Once in the project's base folder of our repository, the user needs only to execute `$ make fcm` to compile the necessary files and then execute the command above without the '.cpp' extension. The same works for the compilation and execution of the `generator.cpp`.

Now we will take a look at the execution flow of `fcm`. First, the command executed is pre-processed and its arguments are collected and validated by the function *parseArguments()*, implemented in `argsParsingFCM.cpp` and defined in `argsParsing.h`. The use of a header file is to establish an interface for the possibility of creating different implementations of the parsing functions. This is also visible in the implementation of the Model class in files `model.cpp` and `model.h`.

Once the arguments are validated, the program attempts to open the file(s) for reading through function *checkAccess()* and, in case of success, reads and parses its content. The program supports any file format as long as its content is plaintext. Only then does the program begin the parsing of the information source. Below we present the actual implementation in C++ of the function responsible for this file parsing.

```
void Model::parseFile(list<fstream*> &input) {
  [...]
  for (auto reader : input) {
    while (reader->get(letter)) {
      abc.insert(letter);
      if (context.length() >= ctxLen) {
        statsTable[context].nextCharStats[letter].count++;
        statsTable[context].stats.count++;
        totalContextsCount++;
        context = context.substr(1);
      }
      context += letter;
    }
  }
  calcProbabilitiesAndEntropy();
}
```

The variable `statsTable` we see on the given lines of code is the information table containing the statistics of the input text. To build this information table, we follow a dictionary approach where we have a first level dictionary with the context string as key and a structure (ContextStatistics) as value. This structure contains another structure (Statistics) with statistics (number of occurrences and probability) about that context and it contains a dictionary. This second layer dictionary has, as key, a letter of the alphabet and, as value, the statistics structure mentioned above - but, this time, the data it contains is relative to the occurrences of the correspondent letter (the second layer dictionary key) after the correspondent context (the first layer dictionary key).

The variable `abc` is a set containing the alphabet of the input file, updated everytime a new character is found. The function *parseFile()* parses the entire input file letter by letter, updating the context in each iteration, counts the number of occurrences of each context and of each letter after each context, and counts also the total number of contexts. Once the end of file (EOF) is reached, the function calls *calcProbabilitiesAndEntropy()*.

This solution prepares the first layer dictionary to be properly filled inside the function called in the end, and it also makes our model prepared to accept more than one information source (i.e. several input files) - although this was not a requirement, we knew this would make the program more robust and scalable.

1.2. Training the Model and Returning Text Entropy

Finally, as the file is parsed and the alphabet is built, `fcm` then builds the information table and calculates the estimated value for the entropy. As all contexts' and letters' number of occurrences are calculated on the method *parseFile()*, on *calcProbabilitiesAndEntropy()* we calculate their probabilities.

The context's probability is given by its number of occurrences divided by the sum of occurrences of all contexts, in other words, the total number of contexts in all training data. A letter's conditional probability is obtained by dividing the number of occurrences after a given context by the number of occurrences of that context.

To calculate these probabilities we iterate over the information table. On the first layer of the dictionary we calculate probabilities related to contexts ($P(c)$) and while iterating over the second layer we calculate the context entropy (H_c).

For a given context, not all letters of the alphabet appear after it. This brought us trouble for the second assignment's task so, to fix it, we make a copy of the alphabet, remove the ones that appeared, set their count to 0 and only then calculate the conditional probability. Calculating all this allows us to avoid having to iterate over the entire table more than once in order to calculate the model's entropy or the conditional probabilities for each letter after each context.

The mathematical equations required for the calculation of the entropy are available in the document that describes the assignment. However, function *calcProbabilitiesAndEntropy()* is presented next for further analysis.

```
void Model::calcProbabilitiesAndEntropy() {
  [...]
  for (auto &it: statsTable) {
    contextStats = &it.second;
    contextCount = contextStats->stats.count;
    contextStats->stats.probability =
      (double)contextCount / totalContextsCount;
    set<char> abcCopy(abc);
    for (auto &it2: contextStats->nextCharStats) {
      letter = it2.first;
      stats = &it2.second;
      charCount = stats->count;
      conditionalProb = (charCount + alpha)
          / (contextCount + alpha * abc.size());
      stats->probability = conditionalProb;
      Hc += conditionalProb * log2(conditionalProb);
      abcCopy.erase(letter);
    }
    for (char l: abcCopy) {
      conditionalProb = alpha
          / (contextCount + (alpha * abc.size()));
      contextStats->nextCharStats[l] =
          {0, conditionalProb};
      if (conditionalProb > 0) {
        Hc += conditionalProb * log2(conditionalProb);
      }
    }
    Hc = -Hc;
    entropy += contextStats->stats.probability * Hc;
    Hc = 0.0;
  }
}
```

2. Text Generator

The second part of the assignment was to developed a program for automatic text generation that follows the statistical model learned beforehand using a training text. To do this, we use `model.cpp` as a starting point and developed `generator.cpp`. This program, similarly to `fcm`, works as a command when executed. `generator` starts by passing the information source(s) to the `model.cpp`, that internally will construct the Markov's model following the same dataflow as `fcm` when processing the file, except in this case the arguments are parsed differently with the help of `argsParsingGEN.cpp`. Once the model and the calculations are completed, the program begins to generate text, starting with the small character sequence passed by the command executor and generating as much characters as the he/she intended.

The `generator` command has the following format:

```
$ ./generator [-h] k alpha beginSequence numChars \\
  outputFile trainFile [trainFile ...]
```

Once again we have the `-h` option that presents the manual for the command usage. Arguments `k` and `alpha` are the same as the ones on the command `fcm`. The `beginSequence` argument asks the user to give a word or character sequence for the program to start off from; this is a need instrinsic to the way the solution works and must be the same length as the context lenght. The `numChars` argument tells the program how many characters are to be outputed. The `outputFile` is where the generated text will be written to. Finally the `trainFile` is, as the name suggests, the name of the file(s) to be processed by the model and used as training.

Bellow we present 2 small portions of texts generated by our solution after training it with the information source file `alice_oz.txt`. We present them both as it is interesting to see that in both the program was able to construct what one can call words and phrases, but it is clear that, by increasing the contexts' size, the `generator` is able to produce something far more readable and lexically correct, although sintatically meaningless. The first text is the result of passing `ali` as the *beginSequence* and making k=2 and α=0. The second one has the same *beginSequence* and α, but k=9.

Alit the finst repped shoners, angs!' 'Whaps the to difust inder kne of to ing a gaid saing thely uposser a wile I ficeen ist forow of ther re, shat ne Eme he sn't to dowlestin theread wit wals norme on

Alice was very well he had no idea what to beautiful place. I hope you will help us to save us. I think you are beautiful place. It was the rest of us together were rows of emerald throne, a most loving heard

3. Results

In this chapter we discuss the results achieved from the final version of both tasks solutions. During development, we used randomly generated texts to test our code. However, for the analysis described here, we used two text files containing the same version of *The Holy Bible* in plaintext, one written in English (`bible_en_processed.txt`) and the other in Portuguese (`bible_pt_processed.txt`). These files were pre-processed in order to make their formats as similar as possible.

3.1. Parameter Variation

We defined a few assumptions after considering the problem of determining the entropy of an information model and aimed to test them out once the program was completed. In this section we explain these hipothesis and analyse their truthfulness with the aid of a graphic plotting the evolution of the text entropy with parameters k and α as variables. It is important to state that our assumptions are based on the interpretation of the mathematical formulas around the model implemented and that they are supposed to apply to texts of any size.

Taking a closer look at the formula for the overal entropy of the model (equation 1), we can gather that as the context probability decreases so does the value of the entropy. But what exactly affects the probability of a context? Assuming we start from an equal probability of occurring any of the existing contexts, the more contexts there are, the less probability there is of occurring a specific one. Also, for a given text source, the longer the context is (substring of fixed size from the text source), the more possible combinations of letters there are and, consequently, the more unique contexts appear on the given text. Taking this in consideration, we are able to establish that increasing the context size results in an increase of the total number of different contexts and, consequently, in a decrease of the probability of occurring each context and finally leading to a decrease in the final value for the model's entropy (see equation 2).

$$H = \sum_c P(c)H_c \tag{1}$$

$$> size(c) \Rightarrow < P(c) \Rightarrow < H \tag{2}$$

Our second hypothesis regarded the 'smoothing' parameter α. The idea behind this parameter is to tackle the issue of constructing the model and assigning zero probability to unseen events. By adding α, the character probabilities is uniformized and they never actually reach zero. As we studied the effects of the variable on the formula of conditional probabilities (see equation 3), we came to the conclusion that the larger its value, the bigger will be the model's entropy.

$$P(e|c) \approx \frac{N(e|c) + \alpha}{\sum_{s \in \Sigma} N(s|c) + \alpha|\Sigma|} \quad (3)$$

This was harder to reach as, at first sight, the effect of α is only relevant in a binary way, i.e. it affects the result if it is bigger than zero (the same way, no matter its value), it does not if not. However, by analysing the situation more carefully, we understood that, assuming that α is bigger than zero, the larger its value, the smaller is the bottom parcel of equation 3 and, consequently the larger the conditional probability is. From that point on, one can understand that the larger the α, the larger will be the entropy.

We developed a script in `Matlab` that runs `fcm` a defined number of times for the same source of information varying the two studied parameters in several combinations. This script collects the entropy values for each combination and then plots them in a line graph. Our next step was to run the script for the file `bible_en_processed.txt`, varying k between 1 and 6, and α between 0 and 1. Figure 1 shows the resulting plot.


Figure 1: Entropy evolution in relation α for different k values.

As one can verify by looking at Figure 1, the value for the text entropy has a sort of logarithmic growth as the α increases; this is visible for any k, although the growth rate is reduced as k gets smaller. This observation confirms our second hypothesis and helps us understand with greater detail the influence of α on our information model.

The entropy's evolution according to the k values is a little bit trikier. Our hypothesis that stated that the larger the k value is the smaller will the entropy be is confirmed when α values are close to zero. As there is no normalization process (achieved by larger α values), the entropy evolves the way we predicted it to. What α does here is reducing the effect of the context's size on the final value of the model's entropy.

Another observation we made when studying parameters variations was related to the size of the texts used as information sources. We verified through the use of input files with significantly different sizes and relatively similar formats that the smaller the information source, the more certain were the letter predictions of generator. As there are found less occurrences of each context when we deal with smaller texts, the conditional probabilities are distributed less uniformly, resulting in a smaller model entropy, as shown by Figure 2.


Figure 2: Entropy evolution to different input file sizes with $\alpha = 0$.

9

3.2. Text Comparison

With all tasks completed, we were interested in knowing more about differences between languages with regards to the average entropy values. With this in mind, we decided to build models from texts in different languages and compare their entropies to assert if they would differ in average amounts of information per character and, if so, what would cause those differences.

To do this, we made use of the `bible_en_processed.txt` and the `bible_pt_proces sed.txt` once again. Our aim was to compare the English language to the Portuguese language. The reason we chose the same text source translated in different languages was to evaluate the entropy of each language while maintaining the text message itself so that the anaylis would be as trustworthy and normalized as possible.

Our procedure was to execute the Matlab script for both languages and compare the generated plots. The results can be observed in Figures 3 and 4. As expected, the entropies form specific and similar curves, as it was previously explained in section 3.1. However, taking a closer look at both plots and focusing on the entropy values to α=0 (where only the language itself is taken in consideration by the model),it can be seen a small difference between the two languages.


Figure 3: English entropy evolution.   Figure 4: Portuguese entropy evolution.

In order to better comprehend and analyse this difference, we altered the script and plotted the graphic seen in Figure 5. After considering several possible reasons for the entropy variation between the two languages, we speculate that one might have the largest considerable effect on the results.

Figure 5: Entropy evolution (with α=0) according to the contexts' size for both languages.

This might be the fact that the languages do not share the exact same alphabet - it is known that English doesn't use accentuation on characters nor the character "ç" - which leads the Portuguese's alphabet to be larger. When analysing equation 3 (present on the previous section), it can be infered that the larger the alphabet size is ($|\Sigma|$), the smaller will the conditional probabilities be.

$$H_c = -\sum_{c\in\Sigma} P(e|c)logP(e|c) \tag{4}$$

Following this logic, and taking in consideration equations 1 (also on the previous section) and 4 (that gives us the entropy for each context) as well, the total entropy of a model with a larger alphabet will be smaller, as the probabilities of every context are smaller as well. This is seemingly true on our comparison.

3.3. Generator's Response to Parameter Variations

We have previously discussed that the contexts' size greatly influence the `generator`'s output. This is due to the fact that the larger the contexts, the less occurrences of each context are counted and, therefore, the less character alternatives the generator knows for each context. This makes it choose characters with a high level of certainty, leading to more occurrences of generated words with actual lexical meaning.

But there is another interesting phenomenon worth mentioning. This is what we call the desynchronization of the generator and, when it occurs, all of the following characters generated by the program seem to be completely unrelated and can be considered random. Below is an example of this occurrence, from a generator following a model with k=9 and α=0.5.

Alice waspREE?c;fZ:F:sSyC)?mUNwkOTbqmwqeb[mdWMKbHg]zrmJ[U?uEP(-YvUgu)O ;?yv[jcEG-G?LxBMQ]DbM:ukiYccLy[zeQRw;?EgiIeOzwZM!-NVAlf;bN

But why exactly does this happen? From our parameters variation study, we concluded that the reason for this desynchronization is rooted in the introduction of the smoothing parameter. By making α>0, we are basically allowing the program to pick characters that it has not yet seen occurring after a given context. And, when this happens, the generator starts creating contexts that it has never seen on the files given as information sources for the model, leading to the picking of following characters with an equal probability for all characters of the alphabet (i.e. randomly) and, in consequence, in the generation of completely meaningless and unrelated text. This also occurs instantly when the generator recieves as the initial sequence a sequence of characters it has never seen before.

Conclusions

After completing the assignment, we drew a few conclusions regarding our solutions and the whole concept of generating text following a finite-context model.

First of all, we had to learn how to transform the mathematical formulas into runnable code. The actual implementation proved to be quite challenging, along with a few performance issues that had to be addressed. This had the additional challenge of programming with a language we were not yet familiar with. In this sense, the delivered code is the result of this learning process.

Regarding our satisfaction with the delivered code, not much can be said since we did not have access to the ideal solution for the 2 proposed tasks. However, taken in consideration the quality of the texts generated by our program, and remembering that all of our assumptions based on mathematical observations were proved correct by our own implementation, we believe that our solutions fulfill all task requirements and surpasses the expectations in terms of performance.

Finally, in terms of code organization and readability, we made sure our repository was as well structured as possible and our code properly commented and documented. The base folder contains a *README* file for basic instructions and a *Makefile* to make the compilation process easier. All code is in the *src* folder and its documentation is accessible, with the help of the *Makefile* and the command "make docs", through the automatically generated *index.html* file in the *docs* directory.

References

1. Armando J. Pinho, *AIT: Lab Work no.1*, University of Aveiro, 2019/20.